

Design a ML system design which predicts in each sentence 'jaguar' refers to car or an animal without LLMs

Example

1. The jaguar is a solitary predator native to the rainforest.
2. I test drove a jaguar F-type at the dealership.
3. Jaguars are often confused with leopards...
4. The Jaguar sedan offers a luxurious ride.

→ Edge Case Example + "jaguar" name ~~of~~ a team

→ Wikipedia, books, wildlife forums, ...
→ Wikipedia, car review, dealership

2) Preprocessing

+ Tokenization

→ The jaguar is a solitary predator native to the rainforest.

['The', 'jaguar', 'is', 'a', 'solitary', 'predator', 'native', 'to', 'the', 'rainforest', '.']

1
[NLTK , Spacy]

- lowercase
- Removing of punctuations & Stop words
(., !, @ ...) (the, as, a, an)
↳ ['jaguar', 'solitary', 'predator', 'native', 'rainforest']

37 Feature engineering

① [TF-IDF] / (Bag of words)

```
from sklearn.feature_extraction.text import TfidfVectorizer
tfidf = TfidfVectorizer(ngram_range=(1, 2))
X_train_tfidf = tfidf.fit_transform(train_sentences)
```

② POS tagging

↳ Identify grammatical roles

The → ' DET '

Jag → ' Noun '

'is' → "Auxiliary"
'solitary' → "adj"] OHE

nom Adv
1 0

③ Dependency parsing

```
The ← det → jaguar
jaguar ← nsubj → is
is ← ROOT → is
a ← det → predator
solitary ← amod → predator
predator ← attr → is
native ← amod → predator
to ← prep → native
the ← det → rainforest
rainforest ← pobj → to
. ← punct → is
```

TF-IDF → sparse numerical representation of word importance

POS tagging → grammatical role of the words

Dependency tree → relationships b/w words

TF-IDF Vector (Vocabulary: [deer, hunts, jaguar]): [0.5, 0.7, 0.3]

POS Tags (One-Hot Encoded): [0, 1, 0, 1, ...] (e.g., 1 for "NOUN" at "jaguar").

Combined Feature Vector: [0.5, 0.7, 0.3, 0, 1, 0, 1, ...]

Approach 1



Modeling

→ Simple model (logistic regression, RF)

Generation of embeddings using modern techniques

1> Static embeddings → GLOVE

2> Contextual embedding → BERT

```
from gensim.models import KeyedVectors
# Load pre-trained GloVe embeddings (50 dimensions)
glove_model = KeyedVectors.load_word2vec_format("glove.6B.50d.txt", binary=False)
# Get embedding for "jaguar"
jaguar_glove = glove_model['jaguar']
print("GloVe Embedding Shape:", jaguar_glove.shape)
print("First 5 dimensions:", jaguar_glove[:5])
```

→ GloVe

```
from transformers import BertTokenizer, BertModel
import torch
tokenizer = BertTokenizer.from_pretrained('bert-base-uncased')
model = BertModel.from_pretrained('bert-base-uncased')
# Example 1: Animal context
sentence_animal = "The jaguar hunts in the rainforest."
inputs_animal = tokenizer(sentence_animal, return_tensors="pt", padding=True, truncation=True)
outputs_animal = model(**inputs_animal)
# Extract embedding for "jaguar" (token position 2)
jaguar_embedding_animal = outputs_animal.last_hidden_state[0, 2, :] # Extract embedding for "jaguar" (token position 2)
# Example 2: Car context
sentence_car = "The Jaguar sedan has a powerful engine."
inputs_car = tokenizer(sentence_car, return_tensors="pt", padding=True, truncation=True)
outputs_car = model(**inputs_car)
# Extract embedding for "Jaguar" (token position 2)
jaguar_embedding_car = outputs_car.last_hidden_state[0, 2, :]
print("Animal Embedding Shape:", jaguar_embedding_animal.shape)
print("Car Embedding Shape:", jaguar_embedding_car.shape)
```

→ BERT


```
similarity = torch.cosine_similarity(jaguar_embedding_animal, jaguar_embedding_car, dim=0)
print(similarity.item()) # Might be ~0.7 (similar but distinct)
```

Deep learning

→ $\begin{bmatrix} 768 \text{ dim} \end{bmatrix}$ → tf. sequence
input dim $[768, 64]$

feed directly

Hybrid approach



Modelling

→ Random forest Classifier] Traditional

→ BiLSTM + Attention

```
import tensorflow as tf
model = tf.keras.Sequential([
    tf.keras.layers.Embedding(input_dim=10000, output_dim=128),
    tf.keras.layers.Bidirectional(tf.keras.layers.LSTM(64, return_sequences=True)),
    tf.keras.layers.Attention(),
    tf.keras.layers.Dense(1, activation="sigmoid")
])
```

Summary ↴

Preprocess Text: Tokenize, clean, remove stopwords.

Extract Features:

TF-IDF, POS tags

BERT/GloVe embeddings

Combine Features: Concatenate or use separately.

Train Model: Logistic Regression, Random Forest, or Neural Network.

```

# Pseudocode
sentences = ["The jaguar hunts...", "The Jaguar sedan..."]
labels = [0, 1]

# Step 4: Traditional Features
tfidf_features = TfidfVectorizer().fit_transform(sentences)
pos_features = OneHotEncoder().fit_transform(extract_pos_tags(sentences))

→ Dependency parsing - # Step 5: Embeddings
bert_embeddings = BertModel().encode(sentences)
glove # Combine
combined_features = concatenate([tfidf_features, pos_features, bert_embeddings])
# Train
model = RandomForestClassifier().fit(combined_features, labels)
top 5%
↓
class_weight: balanced

```

Evaluation

↗ f1-Score
 ↗ Confusion matrix
 { Bert ↗ words BERT focusing on

Deployment

↑ Inference Pipeline

→ format as that of training

Model (Big / Small)

{ optimization → Quantization

{
+ ONNX Runtime → convert model in
ONNX for faster
CPU inference

API deployment

```
from fastapi import FastAPI  
app = FastAPI()  
  
@app.post("/predict")  
async def predict_jaguar(sentence: str):  
    return predict(sentence)
```

Data Collection: Scrape & annotate sentences.


Preprocessing: Tokenize, clean, and filter.

Feature Extraction: Generate BERT embeddings + TF-IDF.

Model Training: Fine-tune BERT or train hybrid model.

Evaluation: Validate with F1-score and error analysis.

Deployment: Optimize and serve via FastAPI.



+ Spacy, NLTK

+ hugging → BERT

+ PyTorch / Tensorflow / Keras

+ Deployment → FastAPI, ONNX, Docker

```
from transformers import BertForSequenceClassification, Trainer, TrainingArguments  
model = BertForSequenceClassification.from_pretrained("bert-base-uncased", num_labels=2)
```

```
training_args = TrainingArguments(  
    output_dir='./results',  
    per_device_train_batch_size=16,  
    num_train_epochs=3,  
    learning_rate=2e-5,  
    evaluation_strategy="epoch",  
)  
trainer = Trainer(  
    model=model,  
    args=training_args,  
    train_dataset=train_dataset, # Hugging Face Dataset object  
    eval_dataset=val_dataset,  
)  
trainer.train()
```

Questions

Imagine you were working on iPhone. Everytime users open their phones, you want to suggest one app they are most likely to open first with 90% accuracy. How would you do that?

An e-commerce company is trying to minimize the time it takes customers to purchase their selected items.
As a machine learning engineer, what can you do to help them?

T-T-P